

Final_Project

2023-03-11

Team: Hengrui Yuan Joanne Teng Jiaheng Dai

#1

```
# data processing
df_train <- read.csv("spam-train.txt")
df_test <- read.csv("spam-test.txt")
for (i in 1:57) {
  names(df_train)[i]=paste("x",i,sep="")
  names(df_test)[i]=paste("x",i,sep="")
}
names(df_train)[58]="Y"
names(df_test)[58]="Y"

# standardization
df_train_std <- df_train
df_train_std[, -58] <- as.data.frame(scale(df_train[, -58]))
df_test_std <- df_test
df_test_std[, -58] <- as.data.frame(scale(df_test[, -58]))

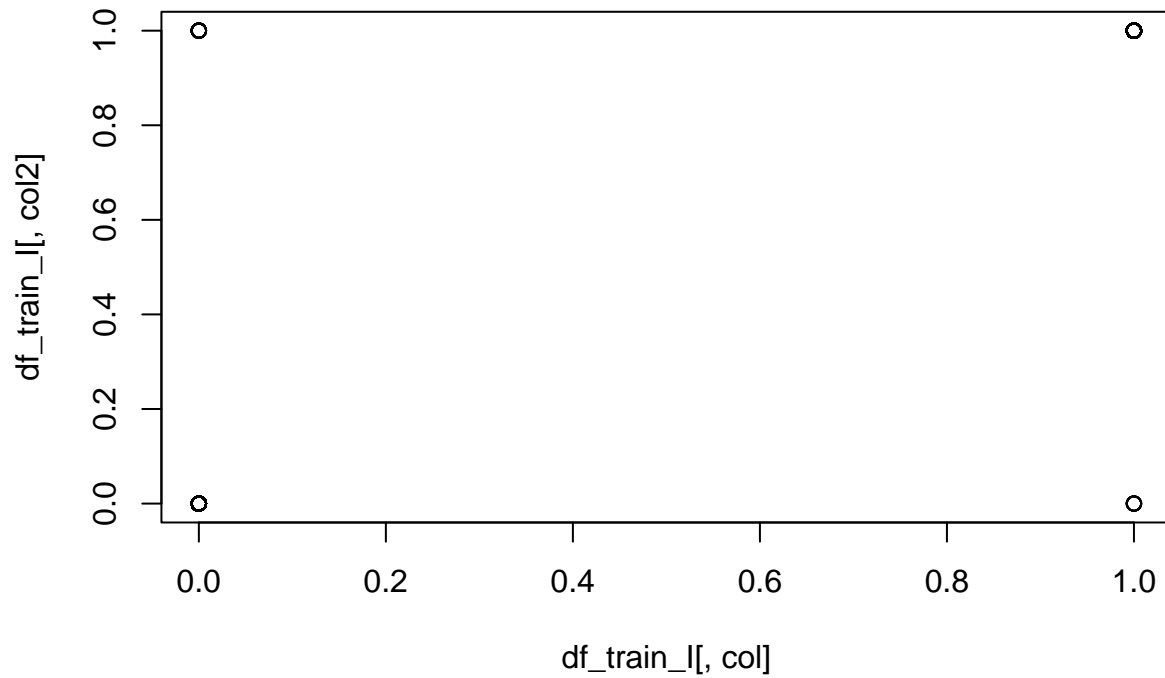
#log transform
df_train_log <- df_train
df_train_log[, -58] <- as.data.frame(log(df_train[, -58]+1))
df_test_log <- df_test
df_test_log[, -58] <- as.data.frame(log(df_test[, -58]+1))

#discretization transform
df_train_I <- as.data.frame(ifelse(df_train>0, 1,0))
df_test_I <- as.data.frame(ifelse(df_test>0, 1,0))
```

#a


```
for (col in 1:ncol(df_train_I)) {
  for (col2 in col:ncol(df_train_I)) {
    if (col!=col2 &&
        !is.na(cor(df_train_I[,col], df_train_I[,col2]))
        && cor(df_train_I[,col], df_train_I[,col2]) > 0.8){
      cat("scatter plot between columns",col,"and",col2,'\n')
      plot(df_train_I[,col], df_train_I[,col2])
    }
  }
}
```

scatter plot between columns 32 and 34




```
for (col in 1:ncol(df_train_std)) {  
  for (col2 in col:ncol(df_train_std)) {  
    if (col!=col2 &&  
        !is.na(cor(df_train_std[,col], df_train_std[,col2]))  
        && cor(df_train_std[,col], df_train_std[,col2]) > 0.75){  
      cat("scatter plot between columns",col,"and",col2,'\n')  
      plot(df_train_std[,col], df_train_std[,col2])  
    }  
  }  
}
```


```
## scatter plot between columns 32 and 34
```



scatter plot between columns 32 and 40

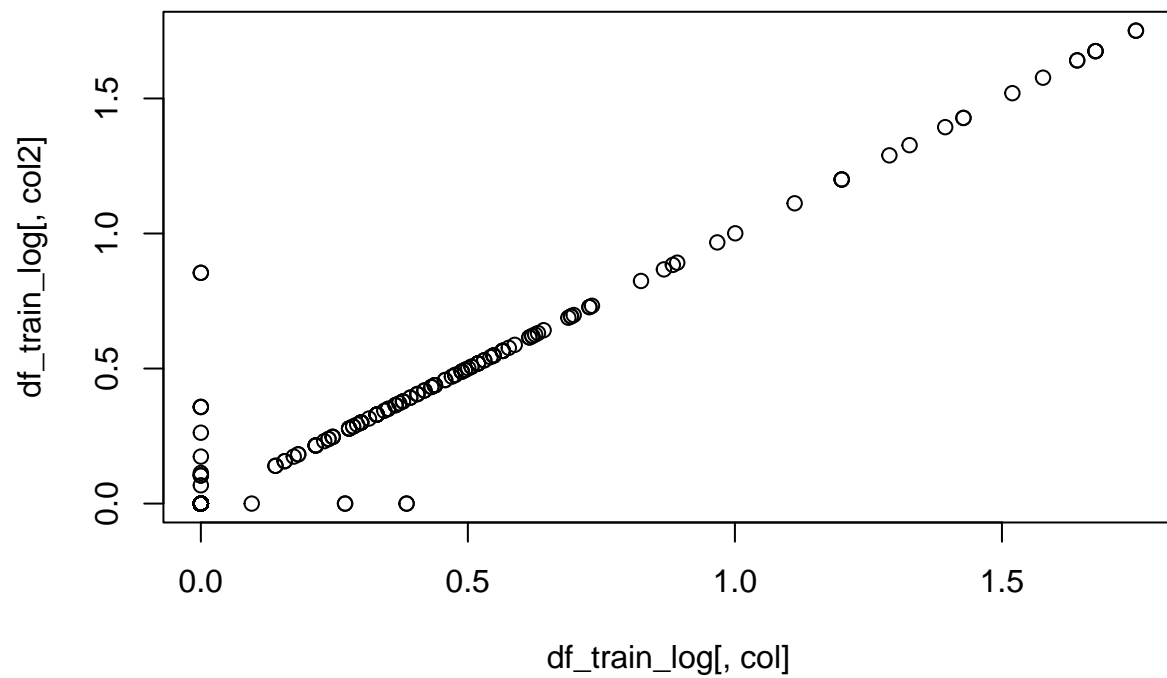


scatter plot between columns 34 and 40




```
for (col in 1:ncol(df_train_log)) {
  for (col2 in col:ncol(df_train_log)) {
    if (col!=col2 &&
        !is.na(cor(df_train_log[,col], df_train_log[,col2]))
        && cor(df_train_log[,col], df_train_log[,col2]) > 0.75){
      cat("scatter plot between columns",col,"and",col2,'\n')
      plot(df_train_log[,col], df_train_log[,col2])
    }
  }
}
```


```
## scatter plot between columns 32 and 34
```




```
## scatter plot between columns 55 and 56
```



```
## scatter plot between columns 56 and 57
```




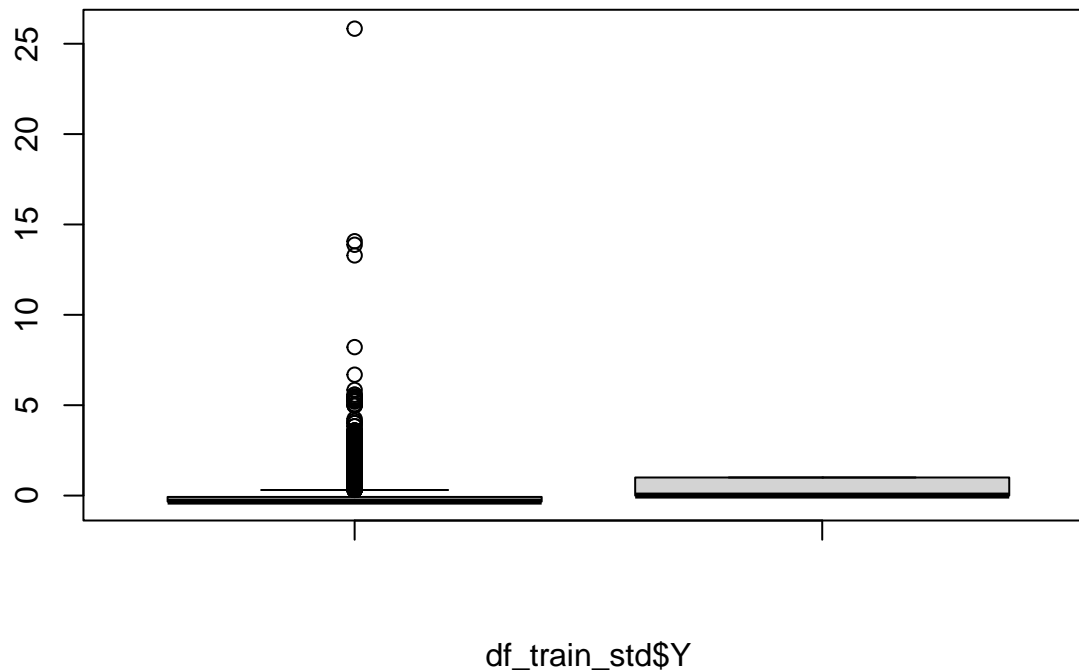
```
# box plot of variables and Y having high correlation
for (col in 1:57){
  if (cor(df_train_log[,col], df_train_log[,58]) > 0.5) {
    boxplot(df_train_log[,col],df_train_log$Y,xlab="df_train_I")
  }
}
```



```
for (col in 1:57){
  if (cor(df_train_std[,col], df_train_std[,58]) > 0.3) {
```

```
boxplot(df_train_std[,col],df_train_std$Y,xlab="df_train_std$Y")
}
```





```
for (col in 1:57){
  if (!is.na(cor(df_train_I[,col], df_train_I[,58])))
    && cor(df_train_I[,col], df_train_I[,58]) > 0.8) {
    boxplot(df_train_I[,col],df_train_I$Y,xlab="df_train_I$Y")
  }
}
```

```
#b
```

```
#logistic regression
#on std train set& test set
Logis_Reg_std_train = glm(as.factor(Y)~., data = df_train_std,
                          family = "binomial")
summary(Logis_Reg_std_train)
```

```
##
## Call:
## glm(formula = as.factor(Y) ~ ., family = "binomial", data = df_train_std)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -4.3245  -0.1988  -0.0001   0.0940   3.6053
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -7.361613   1.762032  -4.178 2.94e-05 ***
## x1            -0.070481   0.085457  -0.825 0.409508
## x2            -0.212713   0.136583  -1.557 0.119379
## x3             0.025730   0.074728   0.344 0.730611
## x4             5.425751   2.634732   2.059 0.039464 *
## x5             0.410339   0.088985   4.611 4.00e-06 ***
```



```

## x6      0.084893  0.057809  1.469 0.141965
## x7      1.307830  0.198298  6.595 4.24e-11 ***
## x8      0.201156  0.073106  2.752 0.005931 **
## x9      0.216452  0.100402  2.156 0.031095 *
## x10     0.057377  0.060904  0.942 0.346146
## x11     -0.195643  0.075244 -2.600 0.009320 **
## x12     -0.035522  0.073025 -0.486 0.626654
## x13     -0.132190  0.110701 -1.194 0.232431
## x14     -0.003391  0.062971 -0.054 0.957058
## x15      0.310890  0.232421  1.338 0.181022
## x16      1.100543  0.164513  6.690 2.24e-11 ***
## x17      0.596501  0.140010  4.260 2.04e-05 ***
## x18     -0.029935  0.083925 -0.357 0.721328
## x19      0.153596  0.077824  1.974 0.048423 *
## x20      1.802268  0.509072  3.540 0.000400 ***
## x21      0.499770  0.085012  5.879 4.13e-09 ***
## x22      0.104744  0.158733  0.660 0.509332
## x23      1.172850  0.241043  4.866 1.14e-06 ***
## x24      0.099468  0.061699  1.612 0.106930
## x25     -3.272100  0.581579 -5.626 1.84e-08 ***
## x26     -0.448614  0.391061 -1.147 0.251310
## x27    -18.555579  3.802484 -4.880 1.06e-06 ***
## x28      0.244576  0.170329  1.436 0.151031
## x29     -2.429257  1.662404 -1.461 0.143936
## x30      0.011451  0.096669  0.118 0.905706
## x31     -0.082969  0.257128 -0.323 0.746941
## x32     -0.374474  0.953628 -0.393 0.694553
## x33     -0.462586  0.246538 -1.876 0.060610 .
## x34      0.853992  1.011821  0.844 0.398661
## x35     -0.611776  0.353254 -1.732 0.083304 .
## x36      0.076193  0.169608  0.449 0.653264
## x37     -0.260527  0.148920 -1.749 0.080214 .
## x38     -0.151496  0.121348 -1.248 0.211871
## x39     -0.026332  0.152992 -0.172 0.863348
## x40     -0.157475  0.176775 -0.891 0.373027
## x41    -18.567059 12.230659 -1.518 0.128995
## x42     -1.695622  0.583196 -2.907 0.003644 **
## x43     -0.454237  0.239231 -1.899 0.057599 .
## x44     -0.734060  0.357167 -2.055 0.039857 *
## x45     -0.885922  0.177301 -4.997 5.83e-07 ***
## x46     -1.085099  0.255167 -4.252 2.11e-05 ***
## x47     -0.642456  0.325245 -1.975 0.048234 *
## x48     -0.502701  0.383351 -1.311 0.189745
## x49     -0.207169  0.101127 -2.049 0.040502 *
## x50      0.047509  0.060039  0.791 0.428765
## x51     -0.065868  0.129006 -0.511 0.609646
## x52      0.248040  0.058142  4.266 1.99e-05 ***
## x53      1.016788  0.162218  6.268 3.66e-10 ***
## x54      0.590673  0.335770  1.759 0.078550 .
## x55     -0.562090  0.229800 -2.446 0.014445 *
## x56      1.082872  0.293775  3.686 0.000228 ***
## x57      0.616630  0.141199  4.367 1.26e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 4120.0 on 3065 degrees of freedom
## Residual deviance: 1157.4 on 3008 degrees of freedom
## AIC: 1273.4
##
## Number of Fisher Scoring iterations: 13

cat("On standardized dataset, x4,x5,x7,x8,x9,x11,x16,x17,x19,x20,x21,x23,x25,x27,
    x42,x44,x45,x46,x47,x49,x52,x53,x55,x56,x57 are statistically significant on
    0.05 significance level.",'\n')

## On standardized dataset, x4,x5,x7,x8,x9,x11,x16,x17,x19,x20,x21,x23,x25,x27,
## x42,x44,x45,x46,x47,x49,x52,x53,x55,x56,x57 are statistically significant on
## 0.05 significance level.

prob_std_train = predict(Logis_Reg_std_train,df_train_std, type='response')
pred_std_train = ifelse(prob_std_train>0.5, "1", "0")
table(pred_std_train, df_train_std$Y) # confusion matrix

##
## pred_std_train    0    1
##                0 1761  133
##                1   87 1085

cat("On standardized train set, we have 133 type one errors; 87 type two errors",'\n')

## On standardized train set, we have 133 type one errors; 87 type two errors

cat("classification error on standardized train set:"
    ,mean(pred_std_train != df_train_std$Y))

## classification error on standardized train set: 0.07175473

prob_std_test = predict(Logis_Reg_std_train,df_test_std, type='response')
pred_std_test = ifelse(prob_std_test>0.5, "1", "0")
table(pred_std_test, df_test_std$Y) # confusion matrix

##
## pred_std_test    0    1
##                0 877  69
##                1  39 548

cat("On standardized test set, we have 69 type one errors; 39 type two errors",'\n')

## On standardized test set, we have 69 type one errors; 39 type two errors
```

```
cat("classification error on standardized test set:"
    ,mean(pred_std_test != df_test_std$Y))
```

```
## classification error on standardized test set: 0.0704501
```

```
#on log train set& test set
Logis_Reg_log_train = glm(as.factor(Y)~., data = df_train_log,
                          family = "binomial")
summary(Logis_Reg_log_train)
```

```
##
## Call:
## glm(formula = as.factor(Y) ~ ., family = "binomial", data = df_train_log)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -4.0831  -0.1647  -0.0010   0.0739   3.7853
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -5.55361    0.47536 -11.683  < 2e-16 ***
## x1            -0.50525    0.52078  -0.970  0.331955
## x2            -0.48375    0.41287  -1.172  0.241326
## x3            -0.34268    0.32461  -1.056  0.291122
## x4             2.49036    2.49963   0.996  0.319109
## x5             1.68052    0.26735   6.286  3.26e-10 ***
## x6             0.49007    0.49976   0.981  0.326780
## x7             3.81919    0.63656   6.000  1.98e-09 ***
## x8             1.11891    0.39254   2.850  0.004366 **
## x9             0.22162    0.61448   0.361  0.718353
## x10            0.20794    0.26664   0.780  0.435468
## x11            -1.73051    0.64790  -2.671  0.007563 **
## x12            -0.13019    0.21705  -0.600  0.548625
## x13            -1.47818    0.59699  -2.476  0.013284 *
## x14             0.49815    0.49244   1.012  0.311725
## x15             2.35454    1.31509   1.790  0.073389 .
## x16             2.00188    0.30550   6.553  5.64e-11 ***
## x17             2.00033    0.49917   4.007  6.14e-05 ***
## x18            -0.62599    0.34041  -1.839  0.065927 .
## x19             0.04967    0.17069   0.291  0.771070
## x20             4.74705    1.75988   2.697  0.006989 **
## x21             0.92793    0.20837   4.453  8.46e-06 ***
## x22             0.19784    0.59582   0.332  0.739858
## x23             3.39784    0.89163   3.811  0.000139 ***
## x24             1.27695    0.41124   3.105  0.001902 **
## x25            -3.97125    0.60153  -6.602  4.06e-11 ***
## x26            -0.43396    0.74531  -0.582  0.560393
## x27            -5.92236    1.42774  -4.148  3.35e-05 ***
## x28             1.27690    0.58913   2.167  0.030201 *
## x29            -5.52545    3.47037  -1.592  0.111344
## x30            -0.08833    0.47636  -0.185  0.852891
## x31            -1.17928    2.44794  -0.482  0.629989
```

```

## x32      -4.26131    4.43663   -0.960  0.336813
## x33      -1.44589    0.73243   -1.974  0.048370 *
## x34       0.86732    4.05415    0.214  0.830598
## x35      -2.60248    1.20496   -2.160  0.030788 *
## x36       0.44061    0.70994    0.621  0.534841
## x37      -1.55260    0.59961   -2.589  0.009615 **
## x38      -1.10219    1.36375   -0.808  0.418971
## x39       0.09939    0.80741    0.123  0.902026
## x40      -1.66153    1.14748   -1.448  0.147621
## x41     -45.30209   35.39196   -1.280  0.200541
## x42      -4.12654    1.24565   -3.313  0.000924 ***
## x43      -5.08561    1.94169   -2.619  0.008815 **
## x44      -2.90440    1.49695   -1.940  0.052354 .
## x45      -2.02986    0.41499   -4.891  1.00e-06 ***
## x46      -2.21581    0.52201   -4.245  2.19e-05 ***
## x47      -7.41902    4.88355   -1.519  0.128716
## x48      -2.02098    1.39842   -1.445  0.148405
## x49      -1.58851    0.79263   -2.004  0.045059 *
## x50      -0.01172    0.62116   -0.019  0.984949
## x51      -3.40427    2.64864   -1.285  0.198691
## x52       2.24783    0.29972    7.500  6.39e-14 ***
## x53       4.93003    0.88667    5.560  2.70e-08 ***
## x54      -0.01276    2.13277   -0.006  0.995225
## x55       0.57047    0.33492    1.703  0.088515 .
## x56       0.09317    0.19497    0.478  0.632742
## x57       0.75138    0.13167    5.707  1.15e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 4120.00  on 3065  degrees of freedom
## Residual deviance:  930.67  on 3008  degrees of freedom
## AIC: 1046.7
##
## Number of Fisher Scoring iterations: 12

cat("On log dataset, x5,x7,x8,x11,x13,,x16,x17,x20,x21,x23,x24,x25,x27,x28,x33,
    x35,x37,x42,x43,x45,x46,x49,x52,x53,x57 are statistically significant on
    0.05 significance level.",'\n')

## On log dataset, x5,x7,x8,x11,x13,,x16,x17,x20,x21,x23,x24,x25,x27,x28,x33,
##    x35,x37,x42,x43,x45,x46,x49,x52,x53,x57 are statistically significant on
##    0.05 significance level.

prob_log_train = predict(Logis_Reg_log_train,df_train_log, type='response')
pred_log_train = ifelse(prob_log_train>0.5, "1", "0")
table(pred_log_train, df_train_log$Y) # confusion matrix

##
## pred_log_train    0    1
##                0 1765   94
##                1   83 1124

```

```
cat("On log train set, we have 94 type one errors; 83 type two errors",'\n')
```

```
## On log train set, we have 94 type one errors; 83 type two errors
```

```
cat("classification error on log train set:"  
    ,mean(pred_log_train != df_train_log$Y))
```

```
## classification error on log train set: 0.05772994
```

```
prob_log_test = predict(Logis_Reg_log_train,df_test_log, type='response')  
pred_log_test = ifelse(prob_log_test>0.5, "1", "0")  
table(pred_log_test, df_test_log$Y) # confusion matrix
```

```
##  
## pred_log_test    0    1  
##               0 879  50  
##               1  37 567
```

```
cat("On log test set, we have 50 type one errors; 37 type two errors",'\n')
```

```
## On log test set, we have 50 type one errors; 37 type two errors
```

```
cat("classification error on log test set:"  
    ,mean(pred_log_test != df_test_log$Y))
```

```
## classification error on log test set: 0.05675147
```

```
#on discretized train set& test set  
Logis_Reg_I_train = glm(as.factor(Y)~., data = df_train_I,  
                        family = "binomial")  
summary(Logis_Reg_I_train)
```

```
##  
## Call:  
## glm(formula = as.factor(Y) ~ ., family = "binomial", data = df_train_I)  
##  
## Deviance Residuals:  
##      Min       1Q   Median       3Q      Max   
## -3.6393  -0.1906  -0.0130   0.0600   3.9295   
##  
## Coefficients: (3 not defined because of singularities)  
##              Estimate Std. Error z value Pr(>|z|)      
## (Intercept) -2.102416   0.189853 -11.074  < 2e-16 ***  
## x1          -0.303288   0.289817  -1.046  0.295340      
## x2          -0.378468   0.275804  -1.372  0.169991      
## x3          -0.199096   0.212662  -0.936  0.349166      
## x4           1.096263   0.824254   1.330  0.183516      
## x5           1.268089   0.216147   5.867  4.44e-09 ***  
## x6           0.251837   0.272999   0.922  0.356278    
```

```

## x7      2.986596  0.386284  7.732 1.06e-14 ***
## x8      0.875956  0.316310  2.769 0.005618 **
## x9      0.228803  0.325213  0.704 0.481713
## x10     0.742334  0.238269  3.116 0.001836 **
## x11    -1.162232  0.334524 -3.474 0.000512 ***
## x12    -0.078386  0.194281 -0.403 0.686606
## x13    -1.161882  0.311431 -3.731 0.000191 ***
## x14     0.941406  0.452030  2.083 0.037286 *
## x15     2.006006  0.693341  2.893 0.003813 **
## x16     1.984574  0.226462  8.763 < 2e-16 ***
## x17     1.096489  0.319793  3.429 0.000606 ***
## x18    -0.857060  0.264974 -3.235 0.001219 **
## x19     0.006173  0.224878  0.027 0.978102
## x20     1.670868  0.554535  3.013 0.002586 **
## x21     0.834541  0.210275  3.969 7.22e-05 ***
## x22     0.811704  0.555362  1.462 0.143857
## x23     1.787932  0.392434  4.556 5.21e-06 ***
## x24     1.385800  0.343260  4.037 5.41e-05 ***
## x25    -3.611805  0.473169 -7.633 2.29e-14 ***
## x26    -0.640909  0.497469 -1.288 0.197628
## x27    -4.432614  0.740639 -5.985 2.17e-09 ***
## x28     1.981100  0.457456  4.331 1.49e-05 ***
## x29    -1.174979  0.668921 -1.757 0.078998 .
## x30    -0.183169  0.519468 -0.353 0.724382
## x31    -1.558354  1.033708 -1.508 0.131673
## x32    -2.211033  1.150862 -1.921 0.054707 .
## x33    -0.926303  0.562100 -1.648 0.099366 .
## x34     0.536590  1.068201  0.502 0.615435
## x35    -0.973396  0.565678 -1.721 0.085295 .
## x36     0.636609  0.417224  1.526 0.127055
## x37    -1.440825  0.348517 -4.134 3.56e-05 ***
## x38     1.173481  0.741366  1.583 0.113453
## x39     0.037746  0.413234  0.091 0.927221
## x40    -0.611575  0.557754 -1.096 0.272861
## x41    -5.823142  3.179724 -1.831 0.067050 .
## x42    -2.410828  0.508740 -4.739 2.15e-06 ***
## x43    -1.500588  0.638112 -2.352 0.018693 *
## x44    -1.301654  0.521225 -2.497 0.012514 *
## x45    -1.391111  0.235936 -5.896 3.72e-09 ***
## x46    -1.789878  0.363562 -4.923 8.52e-07 ***
## x47    -0.695889  1.130614 -0.615 0.538227
## x48    -1.512213  0.617514 -2.449 0.014330 *
## x49    -0.070817  0.275485 -0.257 0.797130
## x50     0.185426  0.196175  0.945 0.344553
## x51    -0.056807  0.409947 -0.139 0.889789
## x52     1.476323  0.186253  7.926 2.26e-15 ***
## x53     1.858610  0.250030  7.434 1.06e-13 ***
## x54    -0.794192  0.338407 -2.347 0.018933 *
## x55      NA      NA      NA      NA
## x56      NA      NA      NA      NA
## x57      NA      NA      NA      NA
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##

```

```

## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 4120.0 on 3065 degrees of freedom
## Residual deviance: 1014.6 on 3011 degrees of freedom
## AIC: 1124.6
##
## Number of Fisher Scoring iterations: 9

cat("On discretized dataset, x5,x7,x8,x10,x11,x13,x14,x15,x16,x17,x18,x20,x21,
    x23,x24,x25,x27,x28,x37,x42,x43,x44,x45,x46,x48,x52,x53,x54 are statistically
    significant on 0.05 significance level.",'\n')

## On discretized dataset, x5,x7,x8,x10,x11,x13,x14,x15,x16,x17,x18,x20,x21,
## x23,x24,x25,x27,x28,x37,x42,x43,x44,x45,x46,x48,x52,x53,x54 are statistically
## significant on 0.05 significance level.

prob_I_train = predict(Logis_Reg_I_train,df_train_I, type='response')
pred_I_train = ifelse(prob_I_train>0.5, "1", "0")
table(pred_I_train, df_train_I$Y) # confusion matrix

##
## pred_I_train    0    1
##              0 1778 105
##              1   70 1113

cat("On discretized train set, we have 105 type one errors; 70 type two errors",'\n')

## On discretized train set, we have 105 type one errors; 70 type two errors

cat("classification error on discretized train set:"
    ,mean(pred_I_train != df_train_I$Y))

## classification error on discretized train set: 0.05707763

prob_I_test = predict(Logis_Reg_I_train,df_test_I, type='response')
pred_I_test = ifelse(prob_I_test>0.5, "1", "0")
table(pred_I_test, df_test_I$Y) # confusion matrix

##
## pred_I_test    0    1
##              0 859  67
##              1  57 550

cat("On discretized test set, we have 67 type one errors; 57 type two errors",'\n') ###

## On discretized test set, we have 67 type one errors; 57 type two errors

```

```
cat("classification error on discretized test set:"
    ,mean(pred_I_test != df_test_I$Y))
```

```
## classification error on discretized test set: 0.08088715
```

```
#c
```

```
#lda & qda
library(MASS)
#lda on standardized dataset
lda_std_train <- lda(as.factor(Y)~.,data=df_train_std)
print(lda_std_train)
```

```
## Call:
## lda(as.factor(Y) ~ ., data = df_train_std)
##
## Prior probabilities of groups:
##      0      1
## 0.6027397 0.3972603
##
## Group means:
##      x1      x2      x3      x4      x5      x6
## 0 -0.1069174  0.03514763 -0.1602290 -0.04479982 -0.2503329 -0.1658362
## 1  0.1622195 -0.05332744  0.2431061  0.06797214  0.3798154  0.2516135
##      x7      x8      x9     x10     x11     x12
## 0 -0.2653823 -0.1561049 -0.1705740 -0.1080780 -0.1928989 -0.009950874
## 1  0.4026491  0.2368489  0.2588019  0.1639804  0.2926742  0.015097878
##      x13     x14     x15     x16     x17     x18     x19
## 0 -0.1170612 -0.03595389 -0.1706871 -0.1798817 -0.1910515 -0.1680778 -0.2368653
## 1  0.1776101  0.05455073  0.2589735  0.2729239  0.2898712  0.2550147  0.3593818
##      x20     x21     x22     x23     x24     x25     x26
## 0 -0.1704232 -0.3148982 -0.07377009 -0.2760402 -0.1857145  0.2142886  0.1932191
## 1  0.2585732  0.4777765  0.11192704  0.4188197  0.2817737 -0.3251276 -0.2931600
##      x27     x28     x29     x30     x31     x32
## 0  0.1424733  0.1392816  0.1081494  0.1270863  0.09765674  0.09291623
## 1 -0.2161665 -0.2113238 -0.1640888 -0.1928206 -0.14816884 -0.14097635
##      x33     x34     x35     x36     x37     x38
## 0  0.0850902  0.08801796  0.1201857  0.1225343  0.1513690  0.02318666
## 1 -0.1291024 -0.13354449 -0.1823507 -0.1859141 -0.2296633 -0.03517976
##      x39     x40     x41     x42     x43     x44
## 0  0.1058538  0.05531977  0.08445125  0.1064889  0.1058379  0.09722614
## 1 -0.1606057 -0.08393344 -0.12813292 -0.1615693 -0.1605817 -0.14751553
##      x45     x46     x47     x48     x49     x50
## 0  0.1320741  0.1258356  0.04790730  0.06393950  0.04390115  0.08091938
## 1 -0.2003884 -0.1909230 -0.07268694 -0.09701165 -0.06660864 -0.12277423
##      x51     x52     x53     x54     x55     x56
## 0  0.05399394 -0.1860095 -0.2738910 -0.02848037 -0.08535606 -0.2423106
## 1 -0.08192184  0.2822214  0.4155588  0.04321159  0.12950575  0.3676437
##      x57
## 0 -0.2335307
## 1  0.3543224
##
```


Coefficients of linear discriminants:

LD1

x1 -0.0526740031
x2 -0.0617536733
x3 0.0733126469
x4 0.0766915865
x5 0.2967065872
x6 0.0794902809
x7 0.3558985084
x8 0.1479128218
x9 0.0295281412
x10 0.0038643527
x11 0.0452143600
x12 -0.0852813073
x13 -0.0048451332
x14 -0.0483942118
x15 -0.0007720204
x16 0.2233197957
x17 0.0833658517
x18 0.1300997822
x19 0.1213140013
x20 0.1137772863
x21 0.3001263541
x22 0.1866456962
x23 0.2671603077
x24 0.1743541760
x25 -0.1713407032
x26 -0.0704307788
x27 -0.1655407020
x28 -0.0276966221
x29 -0.0695662839
x30 -0.0555452458
x31 -0.0164452174
x32 -0.2719340352
x33 -0.0994725729
x34 0.4116799767
x35 -0.0693811328
x36 0.0164011009
x37 -0.0970879515
x38 -0.0602028507
x39 -0.0196065988
x40 -0.0218984182
x41 -0.0085585910
x42 -0.1036432569
x43 -0.0364043829
x44 -0.0983251073
x45 -0.1726127294
x46 -0.1307041959
x47 -0.0520853099
x48 -0.0554336366
x49 -0.1288334500
x50 0.0075783935
x51 -0.0107957126
x52 0.2011100673

```
## x53 0.2654845875
## x54 0.0489521280
## x55 0.0042117602
## x56 0.0721214059
## x57 0.2452356327
```

```
lda_pred_std_train=predict(lda_std_train, df_train_std)
lda_std_train_error=mean(lda_pred_std_train$class!=df_train_std$Y)
cat("Classification error of LDA on standardized train set is",lda_std_train_error,'\n')
```

```
## Classification error of LDA on standardized train set is 0.1017613
```

```
lda_pred_std_test=predict(lda_std_train, df_test_std)
lda_std_test_error=mean(lda_pred_std_test$class!=df_test_std$Y)
cat("Classification error of LDA on standardized test set is",lda_std_test_error,'\n')
```

```
## Classification error of LDA on standardized test set is 0.1030659
```

```
#lda on log dataset
lda_log_train <- lda(as.factor(Y)~.,data=df_train_log)
print(lda_log_train)
```

```
## Call:
## lda(as.factor(Y) ~ ., data = df_train_log)
##
## Prior probabilities of groups:
##      0      1
## 0.6027397 0.3972603
##
## Group means:
##      x1      x2      x3      x4      x5      x6      x7
## 0 0.05315155 0.07959758 0.1375519 0.001253274 0.09927244 0.03851495 0.005018299
## 1 0.12320915 0.11265753 0.2939761 0.025570246 0.35891627 0.14007448 0.198678037
##      x8      x9      x10      x11      x12      x13      x14
## 0 0.02911729 0.03003888 0.1015924 0.01420926 0.2975586 0.04352078 0.02267674
## 1 0.13367902 0.11714435 0.2365405 0.08677934 0.3644886 0.10872460 0.04616668
##      x15      x16      x17      x18      x19      x20      x21
## 0 0.00561555 0.04252057 0.03960108 0.06499006 0.5870133 0.003600182 0.2333573
## 1 0.08059068 0.30381482 0.17613046 0.20519828 1.0562134 0.102872725 0.7281596
##      x22      x23      x24      x25      x26      x27
## 0 0.01117143 0.005257954 0.009195921 0.39115828 0.222766623 0.299920360
## 1 0.06764085 0.160941819 0.152543234 0.01468332 0.009955504 0.001910686
##      x28      x29      x30      x31      x32      x33
## 0 0.112979197 0.0770728349 0.098920500 0.059246231 0.0425018738 0.069848897
## 1 0.007535808 0.0003541024 0.009257066 0.001064792 0.0006326148 0.009409053
##      x34      x35      x36      x37      x38      x39
## 0 0.042327477 0.106264961 0.08920778 0.12410631 0.008446234 0.059000174
## 1 0.002790771 0.008187984 0.01804994 0.02458314 0.003963517 0.009167473
##      x40      x41      x42      x43      x44      x45
## 0 0.04852852 0.0437052621 0.091363412 0.043201950 0.05400213 0.20361459
## 1 0.02872142 0.0001565028 0.002957748 0.008381538 0.00445594 0.08500154
##      x46      x47      x48      x49      x50      x51
```

```

## 0 0.13172703 0.0078437046 0.03439685 0.03016528 0.12674641 0.017786021
## 1 0.01214328 0.0004951255 0.00340204 0.01683878 0.09315826 0.006727094
##          x52          x53          x54          x55          x56          x57
## 0 0.06452462 0.01174497 0.01688401 1.126298 2.388785 4.066896
## 1 0.34984733 0.13587984 0.03553386 1.638094 3.722761 5.393894
##
## Coefficients of linear discriminants:
##          LD1
## x1 -0.575235671
## x2 -0.162912246
## x3 -0.047392824
## x4  0.435354896
## x5  0.715502854
## x6  0.194987045
## x7  1.552542592
## x8  0.657935099
## x9 -0.141445606
## x10 -0.040037396
## x11 0.129640354
## x12 -0.236072919
## x13 -0.394014627
## x14 -0.166042142
## x15 0.069497045
## x16 0.849856603
## x17 0.170920743
## x18 0.253772283
## x19 0.072952107
## x20 0.265146805
## x21 0.481891318
## x22 0.603328790
## x23 1.131377296
## x24 1.181131786
## x25 -0.631451448
## x26 -0.045044927
## x27 -0.236185078
## x28 0.070308337
## x29 -0.303299321
## x30 0.012955773
## x31 0.015237992
## x32 -0.043543551
## x33 -0.587684805
## x34 1.039301922
## x35 -0.434103893
## x36 0.219527939
## x37 -0.446717853
## x38 -0.568521114
## x39 -0.055309147
## x40 -0.311751355
## x41 0.038068320
## x42 -0.389185065
## x43 -0.114871179
## x44 -0.524748092
## x45 -0.388516847
## x46 -0.587599870

```

```
## x47 -1.026554566
## x48 -0.439257727
## x49 -0.996247164
## x50 -0.147432468
## x51 -0.195222568
## x52 1.206070483
## x53 1.729503387
## x54 0.045268764
## x55 -0.001562929
## x56 0.100691654
## x57 0.208945883
```

```
lda_pred_log_train=predict(lda_log_train, df_train_log)
lda_log_train_error=mean(lda_pred_log_train$class!=df_train_log$Y)
cat("Classification error of LDA on log train set is",lda_log_train_error,'\n')
```

```
## Classification error of LDA on log train set is 0.0603392
```

```
lda_pred_log_test=predict(lda_log_train, df_test_log)
lda_log_test_error=mean(lda_pred_log_test$class!=df_test_log$Y)
cat("Classification error of LDA on log test set is",lda_log_test_error,'\n')
```

```
## Classification error of LDA on log test set is 0.1030659
```

```
#qda on standardized dataset
qda_std_train <- qda(as.factor(Y)~.,data=df_train_std)
print(qda_std_train)
```

```
## Call:
## qda(as.factor(Y) ~ ., data = df_train_std)
##
## Prior probabilities of groups:
##      0      1
## 0.6027397 0.3972603
##
## Group means:
##      x1      x2      x3      x4      x5      x6
## 0 -0.1069174 0.03514763 -0.1602290 -0.04479982 -0.2503329 -0.1658362
## 1 0.1622195 -0.05332744 0.2431061 0.06797214 0.3798154 0.2516135
##      x7      x8      x9      x10      x11      x12
## 0 -0.2653823 -0.1561049 -0.1705740 -0.1080780 -0.1928989 -0.009950874
## 1 0.4026491 0.2368489 0.2588019 0.1639804 0.2926742 0.015097878
##      x13      x14      x15      x16      x17      x18      x19
## 0 -0.1170612 -0.03595389 -0.1706871 -0.1798817 -0.1910515 -0.1680778 -0.2368653
## 1 0.1776101 0.05455073 0.2589735 0.2729239 0.2898712 0.2550147 0.3593818
##      x20      x21      x22      x23      x24      x25      x26
## 0 -0.1704232 -0.3148982 -0.07377009 -0.2760402 -0.1857145 0.2142886 0.1932191
## 1 0.2585732 0.4777765 0.11192704 0.4188197 0.2817737 -0.3251276 -0.2931600
##      x27      x28      x29      x30      x31      x32
## 0 0.1424733 0.1392816 0.1081494 0.1270863 0.09765674 0.09291623
## 1 -0.2161665 -0.2113238 -0.1640888 -0.1928206 -0.14816884 -0.14097635
##      x33      x34      x35      x36      x37      x38
```

```
## 0  0.0850902  0.08801796  0.1201857  0.1225343  0.1513690  0.02318666
## 1 -0.1291024 -0.13354449 -0.1823507 -0.1859141 -0.2296633 -0.03517976
##          x39          x40          x41          x42          x43          x44
## 0  0.1058538  0.05531977  0.08445125  0.1064889  0.1058379  0.09722614
## 1 -0.1606057 -0.08393344 -0.12813292 -0.1615693 -0.1605817 -0.14751553
##          x45          x46          x47          x48          x49          x50
## 0  0.1320741  0.1258356  0.04790730  0.06393950  0.04390115  0.08091938
## 1 -0.2003884 -0.1909230 -0.07268694 -0.09701165 -0.06660864 -0.12277423
##          x51          x52          x53          x54          x55          x56
## 0  0.05399394 -0.1860095 -0.2738910 -0.02848037 -0.08535606 -0.2423106
## 1 -0.08192184  0.2822214  0.4155588  0.04321159  0.12950575  0.3676437
##          x57
## 0 -0.2335307
## 1  0.3543224
```

```
qda_pred_std_train=predict(qda_std_train, df_train_std)
qda_std_train_error=mean(qda_pred_std_train$class!=df_train_std$Y)
cat("Classification error of QDA on standardized train set is",qda_std_train_error,'\n')
```

```
## Classification error of QDA on standardized train set is 0.1787345
```

```
qda_pred_std_test=predict(qda_std_train, df_test_std)
qda_std_test_error=mean(qda_pred_std_test$class!=df_test_std$Y)
cat("Classification error of QDA on standardized test set is",qda_std_test_error,'\n')
```

```
## Classification error of QDA on standardized test set is 0.1748206
```

```
#qda on log daataset
qda_log_train <- qda(as.factor(Y)~.,data=df_train_log)
print(qda_log_train)
```

```
## Call:
## qda(as.factor(Y) ~ ., data = df_train_log)
##
## Prior probabilities of groups:
##      0      1
## 0.6027397 0.3972603
##
## Group means:
##          x1          x2          x3          x4          x5          x6          x7
## 0 0.05315155 0.07959758 0.1375519 0.001253274 0.09927244 0.03851495 0.005018299
## 1 0.12320915 0.11265753 0.2939761 0.025570246 0.35891627 0.14007448 0.198678037
##          x8          x9          x10          x11          x12          x13          x14
## 0 0.02911729 0.03003888 0.1015924 0.01420926 0.2975586 0.04352078 0.02267674
## 1 0.13367902 0.11714435 0.2365405 0.08677934 0.3644886 0.10872460 0.04616668
##          x15          x16          x17          x18          x19          x20          x21
## 0 0.00561555 0.04252057 0.03960108 0.06499006 0.5870133 0.003600182 0.2333573
## 1 0.08059068 0.30381482 0.17613046 0.20519828 1.0562134 0.102872725 0.7281596
##          x22          x23          x24          x25          x26          x27
## 0 0.01117143 0.005257954 0.009195921 0.39115828 0.222766623 0.299920360
## 1 0.06764085 0.160941819 0.152543234 0.01468332 0.009955504 0.001910686
##          x28          x29          x30          x31          x32          x33
```

```
## 0 0.112979197 0.0770728349 0.098920500 0.059246231 0.0425018738 0.069848897
## 1 0.007535808 0.0003541024 0.009257066 0.001064792 0.0006326148 0.009409053
##          x34          x35          x36          x37          x38          x39
## 0 0.042327477 0.106264961 0.08920778 0.12410631 0.008446234 0.059000174
## 1 0.002790771 0.008187984 0.01804994 0.02458314 0.003963517 0.009167473
##          x40          x41          x42          x43          x44          x45
## 0 0.04852852 0.0437052621 0.091363412 0.043201950 0.05400213 0.20361459
## 1 0.02872142 0.0001565028 0.002957748 0.008381538 0.00445594 0.08500154
##          x46          x47          x48          x49          x50          x51
## 0 0.13172703 0.0078437046 0.03439685 0.03016528 0.12674641 0.017786021
## 1 0.01214328 0.0004951255 0.00340204 0.01683878 0.09315826 0.006727094
##          x52          x53          x54          x55          x56          x57
## 0 0.06452462 0.01174497 0.01688401 1.126298 2.388785 4.066896
## 1 0.34984733 0.13587984 0.03553386 1.638094 3.722761 5.393894
```

```
qda_pred_log_train=predict(qda_log_train, df_train_log)
qda_log_train_error=mean(qda_pred_log_train$class!=df_train_log$Y)
cat("Classification error of QDA on log train set is",qda_log_train_error,'\n')
```

```
## Classification error of QDA on log train set is 0.1588389
```

```
qda_pred_log_test=predict(qda_log_train, df_test_log)
qda_log_test_error=mean(qda_pred_log_test$class!=df_test_log$Y)
cat("Classification error of QDA on log test set is",qda_log_test_error,'\n')
```

```
## Classification error of QDA on log test set is 0.1572081
```

```
#d
```

```
#linear and nonlinear SVM
library(e1071)
#non linear on std set
tune.gaussian.std=tune(svm, as.factor(Y)~., data=df_train_std,
                      kernel ="radial",
                      ranges=list(cost=c(0.005, 0.05, 0.5),
                                   gamma=c(0.1,0.5,1)))
# Check the selection results
summary(tune.gaussian.std)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost gamma
##   0.5    0.1
##
## - best performance: 0.08415086
##
## - Detailed performance results:
##   cost gamma      error dispersion
```

```
## 1 0.005 0.1 0.39727279 0.03572213
## 2 0.050 0.1 0.27365183 0.03283814
## 3 0.500 0.1 0.08415086 0.01719672
## 4 0.005 0.5 0.39727279 0.03572213
## 5 0.050 0.5 0.39727279 0.03572213
## 6 0.500 0.5 0.13634477 0.01898058
## 7 0.005 1.0 0.39727279 0.03572213
## 8 0.050 1.0 0.39727279 0.03572213
## 9 0.500 1.0 0.21202870 0.03777716
```

```
# Choose the best model (classifier with optimal C)
bestmod.gaussian.std = tune.gaussian.std$best.model
# Check the classifier
summary(bestmod.gaussian.std)
```

```
##
## Call:
## best.tune(METHOD = svm, train.x = as.factor(Y) ~ ., data = df_train_std,
##   ranges = list(cost = c(0.005, 0.05, 0.5), gamma = c(0.1, 0.5,
##   1)), kernel = "radial")
##
##
## Parameters:
##   SVM-Type: C-classification
##   SVM-Kernel: radial
##   cost: 0.5
##
## Number of Support Vectors: 1652
##
## ( 879 773 )
##
##
## Number of Classes: 2
##
## Levels:
## 0 1
```

```
cat("Best cost for gaussian SVM on std dataset is:", 0.5, "gamma:", 0.1, '\n')
```

```
## Best cost for gaussian SVM on std dataset is: 0.5 gamma: 0.1
```

```
gaussian_svm_std_train_error = mean(predict(bestmod.gaussian.std, df_train_std)!=
df_train_std$Y)
gaussian_svm_std_test_error = mean(predict(bestmod.gaussian.std, df_test_std)!=
df_test_std$Y)
cat("gaussian_svm_std_train_error:", gaussian_svm_std_train_error, '\n')
```

```
## gaussian_svm_std_train_error: 0.03065884
```

```
cat("gaussian_svm_std_test_error:", gaussian_svm_std_test_error, '\n')
```

```
## gaussian_svm_std_test_error: 0.07762557
```

```

#linear svm on std set
tune.linear.std=tune(svm, as.factor(Y)~., data=df_train_std,
                     kernel ="linear",
                     ranges=list(cost=c(0.005, 0.05, 0.5)))
# Check the selection results
summary(tune.linear.std)

```

```

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   0.5
##
## - best performance: 0.07306742
##
## - Detailed performance results:
##   cost      error dispersion
## 1 0.005 0.08611271 0.01496371
## 2 0.050 0.07764046 0.01812194
## 3 0.500 0.07306742 0.01456664

```

```

# Choose the best model (classifier with optimal C)
bestmod.linear.std =tune.linear.std$best.model
# Check the classifier
summary(bestmod.linear.std)

```

```

##
## Call:
## best.tune(METHOD = svm, train.x = as.factor(Y) ~ ., data = df_train_std,
##   ranges = list(cost = c(0.005, 0.05, 0.5)), kernel = "linear")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel:  linear
##     cost:  0.5
##
## Number of Support Vectors:  638
##
## ( 329 309 )
##
##
## Number of Classes:  2
##
## Levels:
##  0 1

```

```

cat("Best cost for linear SVM on std dataset is:", 0.5,'\n')

```

```

## Best cost for linear SVM on std dataset is: 0.5

```



```
linear_svm_std_train_error = mean(predict(bestmod.linear.std, df_train_std)!=
                                     df_train_std$Y)
linear_svm_std_test_error = mean(predict(bestmod.linear.std, df_test_std)!=
                                       df_test_std$Y)
cat("linear_svm_std_train_error:", linear_svm_std_train_error, '\n')
```

```
## linear_svm_std_train_error: 0.06621005
```

```
cat("linear_svm_std_test_error:", linear_svm_std_test_error, '\n')
```

```
## linear_svm_std_test_error: 0.06849315
```

```
#non linear on log set
tune.gaussian.log=tune(svm, as.factor(Y)~., data=df_train_log,
                      kernel ="radial",
                      ranges=list(cost=c(0.005, 0.05, 0.5),
                                  gamma=c(0.1,0.5,1)))

# Check the selection results
summary(tune.gaussian.log)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost gamma
##   0.5   0.1
##
## - best performance: 0.08805433
##
## - Detailed performance results:
##   cost gamma      error dispersion
## 1 0.005   0.1 0.39727385 0.01962127
## 2 0.050   0.1 0.34508314 0.01554972
## 3 0.500   0.1 0.08805433 0.01993560
## 4 0.005   0.5 0.39727385 0.01962127
## 5 0.050   0.5 0.39727385 0.01962127
## 6 0.500   0.5 0.21265249 0.01740336
## 7 0.005   1.0 0.39727385 0.01962127
## 8 0.050   1.0 0.39727385 0.01962127
## 9 0.500   1.0 0.22993549 0.01875394
```

```
# Choose the best model (classifier with optimal C)
bestmod.gaussian.log =tune.gaussian.log$best.model
# Check the classifier
summary(bestmod.gaussian.log)
```

```
##
## Call:
## best.tune(METHOD = svm, train.x = as.factor(Y) ~ ., data = df_train_log,
```

```
##      ranges = list(cost = c(0.005, 0.05, 0.5), gamma = c(0.1, 0.5,
##                    1)), kernel = "radial")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##       cost:  0.5
##
## Number of Support Vectors:  1881
##
## ( 1011 870 )
##
##
## Number of Classes:  2
##
## Levels:
##  0 1
```

```
cat("Best cost for gaussian SVM on log dataset is:", 0.5, "gamma:",0.1,'\n')
```

```
## Best cost for gaussian SVM on log dataset is: 0.5 gamma: 0.1
```

```
gaussian_svm_log_train_error = mean(predict(bestmod.gaussian.log, df_train_log)!=
                                     df_train_log$Y)
gaussian_svm_log_test_error = mean(predict(bestmod.gaussian.log, df_test_log)!=
                                         df_test_log$Y)
cat("gaussian_svm_log_train_error:", gaussian_svm_log_train_error, '\n')
```

```
## gaussian_svm_log_train_error: 0.01696021
```

```
cat("gaussian_svm_log_test_error:", gaussian_svm_log_test_error, '\n')
```

```
## gaussian_svm_log_test_error: 0.06066536
```

```
#linear svm on log set
tune.linear.log=tune(svm, as.factor(Y)~., data=df_train_log,
                    kernel ="linear",
                    ranges=list(cost=c(0.005, 0.05, 0.5)))
# Check the selection results
summary(tune.linear.log)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
## 0.005
##
```

```
## - best performance: 0.06067254
##
## - Detailed performance results:
##   cost      error dispersion
## 1 0.005 0.06067254 0.01147780
## 2 0.050 0.06230334 0.01087308
## 3 0.500 0.06426518 0.01294339
```

```
# Choose the best model (classifier with optimal C)
bestmod.linear.log = tune.linear.log$best.model
# Check the classifier
summary(bestmod.linear.log)
```

```
##
## Call:
## best.tune(METHOD = svm, train.x = as.factor(Y) ~ ., data = df_train_log,
##   ranges = list(cost = c(0.005, 0.05, 0.5)), kernel = "linear")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##       cost:  0.005
##
## Number of Support Vectors: 764
##
## ( 383 381 )
##
##
## Number of Classes: 2
##
## Levels:
## 0 1
```

```
cat("Best cost for linear SVM on log dataset is:", 0.005, '\n')
```

```
## Best cost for linear SVM on log dataset is: 0.005
```

```
linear_svm_log_train_error = mean(predict(bestmod.linear.log, df_train_log) !=
                                     df_train_log$Y)
linear_svm_log_test_error = mean(predict(bestmod.linear.log, df_test_log) !=
                                       df_test_log$Y)
cat("linear_svm_log_train_error:", linear_svm_log_train_error, '\n')
```

```
## linear_svm_log_train_error: 0.05936073
```

```
cat("linear_svm_log_test_error:", linear_svm_log_test_error, '\n')
```

```
## linear_svm_log_test_error: 0.06066536
```

```

#non linear on I set
tune.gaussian.I=tune(svm, as.factor(Y)~., data=df_train_I,
                      kernel ="radial",
                      ranges=list(cost=c(0.005, 0.05, 0.5),
                                   gamma=c(0.1,0.5,1)))
# Check the selection results
summary(tune.gaussian.I)

```

```

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost gamma
##   0.5   0.5
##
## - best performance: 0.04859701
##
## - Detailed performance results:
##   cost gamma   error dispersion
## 1 0.005   0.1 0.30561623 0.02138926
## 2 0.050   0.1 0.07926167 0.01707658
## 3 0.500   0.1 0.05675523 0.01548598
## 4 0.005   0.5 0.39726320 0.02132422
## 5 0.050   0.5 0.33430521 0.02429868
## 6 0.500   0.5 0.04859701 0.01361423
## 7 0.005   1.0 0.39726320 0.02132422
## 8 0.050   1.0 0.39726320 0.02132422
## 9 0.500   1.0 0.12131741 0.01754943

```

```

# Choose the best model (classifier with optimal C)
bestmod.gaussian.I =tune.gaussian.I$best.model
# Check the classifier
summary(bestmod.gaussian.I)

```

```

##
## Call:
## best.tune(METHOD = svm, train.x = as.factor(Y) ~ ., data = df_train_I,
##   ranges = list(cost = c(0.005, 0.05, 0.5), gamma = c(0.1, 0.5,
##   1)), kernel = "radial")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##   cost:  0.5
##
## Number of Support Vectors:  1706
##
## ( 885 821 )
##

```

```

##
## Number of Classes:  2
##
## Levels:
##  0 1

cat("Best cost for gaussian SVM on discretized dataset is:", 0.5, "gamma:",0.5,'\n')

## Best cost for gaussian SVM on discretized dataset is: 0.5 gamma: 0.5

gaussian_svm_I_train_error = mean(predict(bestmod.gaussian.I, df_train_I)!=
                                     df_train_I$Y)
gaussian_svm_I_test_error  = mean(predict(bestmod.gaussian.I, df_test_I)!=
                                     df_test_I$Y)
cat("gaussian_svm_I_train_error:", gaussian_svm_I_train_error, '\n')

## gaussian_svm_I_train_error: 0.02217873

cat("gaussian_svm_I_test_error:", gaussian_svm_I_test_error, '\n')

## gaussian_svm_I_test_error: 0.0541422

#linear svm on I set
tune.linear.I=tune(svm, as.factor(Y)~., data=df_train_I,
                  kernel ="linear",
                  ranges=list(cost=c(0.005, 0.05, 0.5)))
# Check the selection results
summary(tune.linear.I)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   0.5
##
## - best performance: 0.06816121
##
## - Detailed performance results:
##   cost      error dispersion
## 1 0.005 0.08578272 0.01317338
## 2 0.050 0.07077452 0.01278327
## 3 0.500 0.06816121 0.01815947

# Choose the best model (classifier with optimal C)
bestmod.linear.I =tune.linear.I$best.model
# Check the classifier
summary(bestmod.linear.I)

```

```
##
## Call:
## best.tune(METHOD = svm, train.x = as.factor(Y) ~ ., data = df_train_I,
##           ranges = list(cost = c(0.005, 0.05, 0.5)), kernel = "linear")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##           cost: 0.5
##
## Number of Support Vectors: 582
##
## ( 294 288 )
##
##
## Number of Classes: 2
##
## Levels:
## 0 1
```

```
cat("Best cost for linear SVM on discretized dataset is:", 0.5, '\n')
```

```
## Best cost for linear SVM on discretized dataset is: 0.5
```

```
linear_svm_I_train_error = mean(predict(bestmod.linear.I, df_train_I)!=
                                   df_train_I$Y)
linear_svm_I_test_error = mean(predict(bestmod.linear.I, df_test_I)!=
                                      df_test_I$Y)
cat("linear_svm_I_train_error:", linear_svm_I_train_error, '\n')
```

```
## linear_svm_I_train_error: 0.06001305
```

```
cat("linear_svm_I_test_error:", linear_svm_I_test_error, '\n')
```

```
## linear_svm_I_test_error: 0.07371168
```

```
#e
```

```
library(tree)
library(randomForest)
tree_std <- tree(as.factor(Y)~., data = df_train_std)
#cv pruned tree
cross_val = cv.tree(tree_std, K=10)
cv_size = cross_val$size[which.min(cross_val$dev)]
tree_prune_std = prune.tree(tree_std, best=cv_size)
summary(tree_prune_std)
```

```
##
## Classification tree:
## tree(formula = as.factor(Y) ~ ., data = df_train_std)
```

```

## Variables actually used in tree construction:
## [1] "x52" "x7" "x24" "x16" "x23" "x27" "x53" "x55" "x25"
## Number of terminal nodes: 13
## Residual mean deviance: 0.4929 = 1505 / 3053
## Misclassification error rate: 0.08774 = 269 / 3066

ptree_std_train_error=mean(predict(tree_std, df_train_std, type="class")!=df_train_std$Y)
ptree_std_test_error=mean(predict(tree_std, df_test_std, type="class")!=df_test_std$Y)
#bagging
bag_std = randomForest(as.factor(Y)~., data = df_train_std,
                        mtry=(ncol(df_train_std)-1), importance=TRUE, ntree=100)
bag_std_train_error=mean(predict(bag_std, df_train_std, type="class")!=df_train_std$Y)
bag_std_test_error=mean(predict(bag_std, df_test_std, type="class")!=df_test_std$Y)
cat("error of pruned tree on std train set: ", ptree_std_train_error, '\n')

## error of pruned tree on std train set: 0.08773646

cat("error of pruned tree on std test set: ", ptree_std_test_error, '\n')

## error of pruned tree on std test set: 0.1272016

cat("error of tree with bagging on std train set: ", bag_std_train_error, '\n')

## error of tree with bagging on std train set: 0.0003261579

cat("error of tree with bagging on std test set: ", bag_std_test_error, '\n')

## error of tree with bagging on std test set: 0.04631442

tree_log <- tree(as.factor(Y)~., data = df_train_log)
#cv pruned tree
cross_val =cv.tree(tree_log, K=10)
cv_size = cross_val$size[which.min(cross_val$dev)]
tree_prune_log = prune.tree(tree_log, best=cv_size)
summary(tree_prune_log)

##
## Classification tree:
## tree(formula = as.factor(Y) ~ ., data = df_train_log)
## Variables actually used in tree construction:
## [1] "x52" "x7" "x24" "x16" "x23" "x27" "x53" "x55" "x25"
## Number of terminal nodes: 13
## Residual mean deviance: 0.4929 = 1505 / 3053
## Misclassification error rate: 0.08774 = 269 / 3066

ptree_log_train_error=mean(predict(tree_log, df_train_log, type="class")!=df_train_log$Y)
ptree_log_test_error=mean(predict(tree_log, df_test_log, type="class")!=df_test_log$Y)
#bagging
bag_log = randomForest(as.factor(Y)~., data = df_train_log,
                        mtry=(ncol(df_train_log)-1), importance=TRUE, ntree=100)
bag_log_train_error=mean(predict(bag_log, df_train_log, type="class")!=df_train_log$Y)
bag_log_test_error=mean(predict(bag_log, df_test_log, type="class")!=df_test_log$Y)
cat("error of pruned tree on log train set: ", ptree_log_train_error, '\n')

```

```
## error of pruned tree on log train set: 0.08773646
```

```
cat("error of pruned tree on log test set: ", ptree_log_test_error, '\n')
```

```
## error of pruned tree on log test set: 0.09067189
```

```
cat("error of tree with bagging on log train set: ", bag_log_train_error, '\n')
```

```
## error of tree with bagging on log train set: 0
```

```
cat("error of tree with bagging on log test set: ", bag_log_test_error, '\n')
```

```
## error of tree with bagging on log test set: 0.03652968
```

```
tree_I <- tree(as.factor(Y)~., data = df_train_I)
#cv pruned tree
cross_val =cv.tree(tree_I, K=10)
cv_size = cross_val$size[which.min(cross_val$dev)]
tree_prune_I = prune.tree(tree_I, best=cv_size)
summary(tree_prune_I)
```

```
##
## Classification tree:
## tree(formula = as.factor(Y) ~ ., data = df_train_I)
## Variables actually used in tree construction:
## [1] "x7" "x53" "x16" "x25" "x52" "x45"
## Number of terminal nodes: 8
## Residual mean deviance: 0.5656 = 1730 / 3058
## Misclassification error rate: 0.1171 = 359 / 3066
```

```
ptree_I_train_error=mean(predict(tree_I, df_train_I, type="class")!=df_train_I$Y)
ptree_I_test_error=mean(predict(tree_I, df_test_I, type="class")!=df_test_I$Y)
#bagging
bag_I = randomForest(as.factor(Y)~., data = df_train_I,
                      mtry=(ncol(df_train_I)-1), importance=TRUE, ntree=100)
bag_I_train_error=mean(predict(bag_I, df_train_I, type="class")!=df_train_I$Y)
bag_I_test_error=mean(predict(bag_I, df_test_I, type="class")!=df_test_I$Y)
cat("error of pruned tree on discretized train set: ", ptree_I_train_error, '\n')
```

```
## error of pruned tree on discretized train set: 0.1170907
```

```
cat("error of pruned tree on discretized test set: ", ptree_I_test_error, '\n')
```

```
## error of pruned tree on discretized test set: 0.1226354
```

```
cat("error of tree with bagging on discretized train set: ", bag_I_train_error, '\n')
```

```
## error of tree with bagging on discretized train set: 0.006849315
```



```
cat("error of tree with bagging on discretized test set: ", bag_I_test_error, '\n')
```

```
## error of tree with bagging on discretized test set: 0.04631442
```

```
#tabulate classification errors for different models and dataset
```

```
library("knitr")
err_T <- matrix(nrow = 7, ncol = 6)
colnames(err_T) <- c("Std_train", "Std_test", "Log_train", "Log_test",
                    "Discretized_train", "Discretized_test")
row.names(err_T) <- c("Logistic_Regression", "LDA", "QDA", "Linear_SVM",
                    "Gaussian_SVM", "Pruned_Tree", "Bagging")

err_T[1,1] = mean(pred_std_train != df_train_std$Y)
err_T[1,2] = mean(pred_std_test != df_test_std$Y)
err_T[1,3] = mean(pred_log_train != df_train_log$Y)
err_T[1,4] = mean(pred_log_test != df_test_log$Y)
err_T[1,5] = mean(pred_I_train != df_train_I$Y)
err_T[1,6] = mean(pred_I_test != df_test_I$Y)

err_T[2,1] = lda_std_train_error
err_T[2,2] = lda_std_test_error
err_T[2,3] = lda_log_train_error
err_T[2,4] = lda_log_test_error
err_T[2,5] = NA
err_T[2,6] = NA

err_T[3,1] = qda_std_train_error
err_T[3,2] = qda_std_test_error
err_T[3,3] = qda_log_train_error
err_T[3,4] = qda_log_test_error
err_T[3,5] = NA
err_T[3,6] = NA

err_T[4,1] = linear_svm_std_train_error
err_T[4,2] = linear_svm_std_test_error
err_T[4,3] = linear_svm_log_train_error
err_T[4,4] = linear_svm_log_test_error
err_T[4,5] = linear_svm_I_train_error
err_T[4,6] = linear_svm_I_test_error

err_T[5,1] = gaussian_svm_std_train_error
err_T[5,2] = gaussian_svm_std_test_error
err_T[5,3] = gaussian_svm_log_train_error
err_T[5,4] = gaussian_svm_log_test_error
err_T[5,5] = gaussian_svm_I_train_error
err_T[5,6] = gaussian_svm_I_test_error

err_T[6,1] = ptree_std_train_error
err_T[6,2] = ptree_std_test_error
err_T[6,3] = ptree_log_train_error
err_T[6,4] = ptree_log_test_error
err_T[6,5] = ptree_I_train_error
```

```

err_T[6,6] = ptree_I_test_error

err_T[7,1] = bag_std_train_error
err_T[7,2] = bag_std_test_error
err_T[7,3] = bag_log_train_error
err_T[7,4] = bag_log_test_error
err_T[7,5] = bag_I_train_error
err_T[7,6] = bag_I_test_error
kable(err_T)

```

	Std_train	Std_test	Log_train	Log_test	Discretized_train	Discretized_test
Logistic_Regression	0.0717547	0.0704501	0.0577299	0.0567515	0.0570776	0.0808871
LDA	0.1017613	0.1030659	0.0603392	0.0652316	NA	NA
QDA	0.1787345	0.1748206	0.1588389	0.1572081	NA	NA
Linear_SVM	0.0662100	0.0684932	0.0593607	0.0606654	0.0600130	0.0737117
Gaussian_SVM	0.0306588	0.0776256	0.0169602	0.0606654	0.0221787	0.0541422
Pruned_Tree	0.0877365	0.1272016	0.0877365	0.0906719	0.1170907	0.1226354
Bagging	0.0003262	0.0463144	0.0000000	0.0365297	0.0068493	0.0463144

Based on the Table of Classification errors, comparatively, tree based classifiers with bagging have best overall performance on three versions of transformed data sets.

In conclusion, we recommend performing the classification task with tree based classifiers with bagging because the classification errors on three versions of transformed data set are smallest among all models applied.